

Recognition

Allows us to see what is in an image - used in tracking, identification, reading characters, controlling robots etc...

- Template Matching

- Applications

- Searching - locating known objects
- Recognition - e.g. reading number plates (not a good approach)
- Visual Inspection - e.g. golden template matching
 - Used in manufacturing: the perfect item is used as a template, compared to the items being produced, and a difference image is examined to detect flaws
- Stereo Vision - determining how far away an object is using small templates
- Tracking - assume an object changes slightly frame to frame, can use small templates to maintain the tracking

- Matching Algorithm

- Basic algorithm:
 - Input image and object
 - Evaluate a match criterion for each possible position of the object in the image
 - Search for local maxima above a threshold
- Problems
 - To consider every possible position, do we evaluate at every rotation and scale also?
 - How do we identify a threshold? Too high and you miss genuine matches, too low you introduce false positives
 - What is the match criteria?

- Matching Criteria (metrics)

$$D_{\text{SquareDifferences}}(i, j) = \sum_{(m,n)} (f(i + m, j + n) - t(m, n))^2$$

$$D_{\text{NormalisedSquareDifferences}}(i, j) = \frac{\sum_{(m,n)} (f(i + m, j + n) - t(m, n))^2}{\sqrt{\sum_{(m,n)} f(i + m, j + n)^2 \sum_{(m,n)} t(m, n)^2}}$$

$$D_{\text{CrossCorrelation}}(i, j) = \sum_{(m,n)} f(i + m, j + n) \cdot t(m, n)$$

$$D_{\text{NormalisedCrossCorrelation}}(i, j) = \frac{\sum_{(m,n)} f(i + m, j + n) \cdot t(m, n)}{\sqrt{\sum_{(m,n)} f(i + m, j + n)^2 \sum_{(m,n)} t(m, n)^2}}$$

- - **Normalised square differences: 0 is optimal, 1 is the worst**
 - **Normalised cross correlation: 0 is the worst, 1 is optimal**
- **Cross correlation ignores luminance**
- What to do at boundaries?

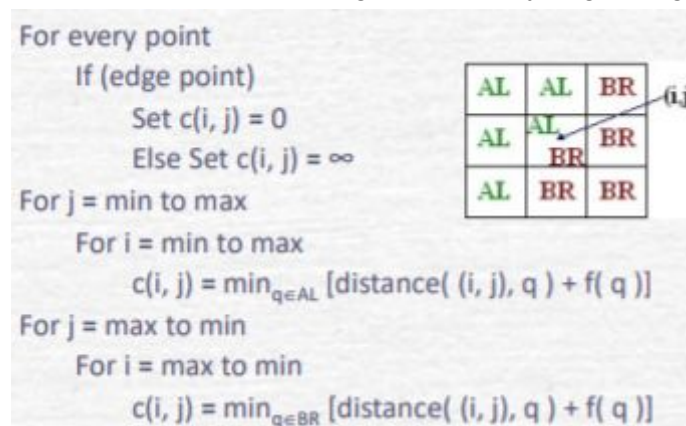
- The degree-of-fit will be inconsistent with neighbouring values, so don't compute generally

- Control Strategies

- Localise close copies
 - Similar values in close proximity are probably matches to the same object so we select the greatest value in a local neighbourhood
 - The larger the neighbourhood the more dispersed the matches (typically 8)
- Process using image hierarchy
 - **Use low res first, and limit the regions for high res matching based off this**
 - Not appropriate for very detailed templates
- **Prioritise searches using the low res image matches as probability**

- Chamfer Matching

- Template matching requires very close matches
 - Generally not possible
 - **Noise, Orientation**
- **Need to be more flexible**
- **A chamfer image is an encoding for each pixel of its distance to the nearest object (edge)**
- Compute a chamfered image for a binary edge image




- 2 pass algorithm
 - First pass considers all points AL (above and left) of the current pixel
 - Second pass considers all points BR (below and right) of the current pixel
- **The template in Chamfer matching is a binary template in which only the object pixels are set and considered**
 - The metric is the sum of the overlapping values (0 is optimal)

- Statistical Pattern Recognition (SPR)

- Derive features from known objects, and use these features to classify unknowns based on their similarity to these shapes

- **Probability**
 - $P(A) = N(A)/n$ as n tends to infinity
 - **Independent events:** $P(AB) = P(A)P(B)$
 - **Dependent events:** $P(AB) = P(A|B)P(B)$
- Typical Probability Problem in this context
 - Having identified feature value x , given a class W_i , what's the probability that the object is of the class W_i ? $P(x|W_i)$
 - This is in effect a **probability density function of the likelihood of a value occurring for that class**
 - We compute these for multiple classes
 - However $P(x|W_i)$ is a priori probability, it is reverse engineering (likelihood of a value being in this class, given this class), we care more about the probability of the class given some value [$P(W_i|x)$]
 - **Bayes Theorem**
 - For two classes A and B, the a-posteriori probability is: $P(B|A) = P(A|B)P(B)/P(A)$
 - Where W_i forms a partitioning of the event space:

$$\frac{p(x | W_i)P(W_i)}{\sum_j p(x | W_j)P(W_j)}$$



- - Example probability density functions for two classes on a single feature x
 - The classes aren't linearly separable (as they overlap) - need another feature
- **Example features:** Area, MinboundingRect, Convex hull(smallest convex region) , elongatedness (area divided by thickness squared), concavities and holes, perimeter length and circularity

- **Advantages of SPR: accurate**
- **Disadvantage: training, requires a decently sized training set**
- Learning strategies:
 - Supervised: probability density estimation
 - Training set includes class spec for every instance
 - Unsupervised: Cluster analysis
 - Look for similarities in feature space

- Support Vector Machines (SVM)

- Works for 2 class problem
- Consider 2 linearly separable classes in n-dimensional feature space
- **Finds the optimal hyperplane (max margin between classes)**
- Given many n-dimensional training samples x_i with associated class identifiers $w_i = \{-1, 1\}$
 - Define separating hyperplanes
 $w.x + b = 0$, $w.x + b = 1$, $w.x + b = -1$
 - Constraint: $w_i(w.x_i + b) \geq 1$
 - If x^+ is a point on $w.x+b=1$ and x^- is the nearest point on $w.x+b=-1$
 - Then $\lambda w = x^+ - x^-$
 - Multiply by w ...
 - $\lambda w.w = (1 - b) - (-1 - b) = 2$
 - Therefore $\lambda w.w = \lambda \|w\|^2 = 2$ as $\|w\| = \sqrt{w.w}$

$$\|x^+ - x^-\| = \|\lambda w\| = \left\| \frac{2.w}{\|w\|^2} \right\| = \left\| \frac{2.w}{w.w} \right\| = \left\| \frac{2}{w} \right\| = \frac{2}{\|w\|}$$

- Hence we must minimise $\|w\|$ to maximise the separation, subject to
 $w_i(w.x_i + b) \geq 1$

- Lagrangian Optimisation

- This is a minimisation problem subject to constraints which is amenable to solution by Lagrangian optimisation

$$L(w, b, \alpha) = \frac{\|w\|^2}{2} - \sum_{i=1}^n \alpha_i [\omega_i \cdot (w.x_i + b) - 1]$$

- To find the minimums we set the partial differentials to 0
 - Classification based on:
 - Entering an unknown feature vector
 - Into the original equation $f(x_i) = w.x_i + b$
 - And classifying based on whether the value is positive or negative
- What if classes aren't linearly separable?
 - Apply the kernel trick
 - Dot products are replaced by some non-linear kernel function
 - Soft margin or multiclass SVM

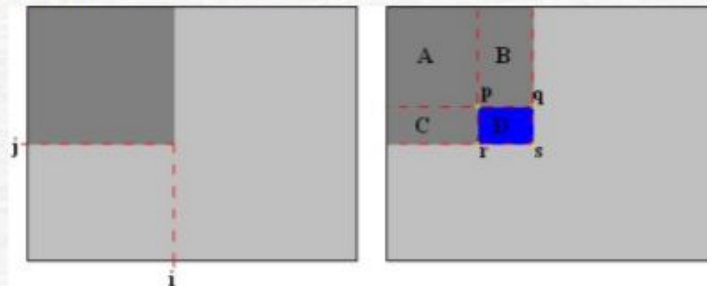
- **Cascade of Haar classifiers (Haar)**

- Robust object detection using a cascade of classifiers
- **Efficient calculation of features**
 - Only uses simple features
- Learns based on a number of positive and negative samples
- Selects a large number of features during training and creates classifiers with them to accept/reject
 - Classifiers are organised into a cascade(sequentially) where if the sub image is rejected by any of the classifiers, it is rejected
 - **Efficient: most sub-images stopped by first or second classifier**
- System trains using particular scales, but classifiers are designed so that they are easily resized
- **Features are determined as the difference of the sums of a number of rectangular regions**
 - Place the mask in a specific location
 - At a specific scale
 - Subtract the normalized sum of the white pixels from the normalized sum of the black pixels
 - Why does this work?
 - **Integral image**

Every point $ii(i,j) = \sum_{i'=0..i} \sum_{j'=0..j} \text{image}(i',j')$

Sum of points in rectangle D:

$$\text{sum}(D) = ii(p) + ii(s) - ii(q) - ii(r)$$



- **Features can be computed at any scale for the same cost**

- **Training**

- Needs a large number of positive and negative samples
- Calculates hundreds of thousands of possible features
- Uses training to identify which features are most important and at which points in the cascade

- **Weak Classifiers**

- Created by combining a specific feature with a threshold, making an accept or reject classification through comparison

$$p_j \text{feature}_j(x) < p_j \vartheta_j$$

Tune threshold (ϑ_j)

- **Strong Classifiers**

- Combine a number of weak classifiers
- Algorithm for this combination is AdaBoost

- **Principal Components Analysis (PCA)**

- Statistical Technique for: Data Analysis, Compression, and Recognition
- Analyses data covariance
- Identifies principal directions
- **Example**
 - Given 2D data, and N samples/vectors
 - Find the mean sample, the direction of maximum covariance, and the direction orthogonal to this

- **Background - Eigenvalues and Eigenvectors**

- Consider a square matrix A
- The eigenvalues of A are the roots of the characteristic equation:

- $\text{determinant}(A - \lambda I) = |A - \lambda I| = 0$

- **For each λ there will be an eigenvector x such that $Ax = \lambda x$**
- An $n \times n$ matrix will have n eigenvalues
- Consider $n = 2$; There will be 2 eigenvalues: λ_1, λ_2
 - With corresponding eigenvectors: x_1, x_2
 - Where $Ax_1 = \lambda_1 x_1$ and $Ax_2 = \lambda_2 x_2$

Combining we get $A \begin{bmatrix} x_1 & x_2 \end{bmatrix} = \begin{bmatrix} x_1 & x_2 \end{bmatrix} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix}$

and $A\Phi = \Phi\Lambda$

where $\Phi = \begin{bmatrix} x_1 & x_2 \end{bmatrix}$ and $\Lambda = \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix}$

If we normalise the eigenvectors then $\Phi\Phi^T = \Phi^T\Phi = I$

so $\Phi^T A \Phi = \Phi^T \Phi \Lambda = \Lambda$

and $A = \Lambda \Phi \Phi^T = \Phi \Lambda \Phi^T$

- In PCA we have N samples/vectors in some n-dimensional space
- Combine into data matrix D

$$D = \begin{bmatrix} x_1^1 & x_2^1 & \dots & x_n^1 \\ x_1^2 & x_2^2 & \dots & x_n^2 \\ \vdots & \vdots & \ddots & \vdots \\ x_1^N & x_2^N & \dots & x_n^N \end{bmatrix}$$

- Each row is a sample
- **Determine the mean of the samples**

- $\mu = \left[\sum_{i=1}^N x_1^i / N \quad \sum_{i=1}^N x_2^i / N \quad \dots \quad \sum_{i=1}^N x_n^i / N \right]$

- **Compute mean-centred data U**

$$U = D - \begin{bmatrix} \mu \\ \mu \\ \vdots \\ \mu \end{bmatrix} = \begin{bmatrix} x_1^1 - \sum_{i=1}^N x_1^i / N & x_2^1 - \sum_{i=1}^N x_2^i / N & \dots & x_n^1 - \sum_{i=1}^N x_n^i / N \\ x_1^2 - \sum_{i=1}^N x_1^i / N & x_2^2 - \sum_{i=1}^N x_2^i / N & \dots & x_n^2 - \sum_{i=1}^N x_n^i / N \\ \vdots & \vdots & \ddots & \vdots \\ x_1^N - \sum_{i=1}^N x_1^i / N & x_2^N - \sum_{i=1}^N x_2^i / N & \dots & x_n^N - \sum_{i=1}^N x_n^i / N \end{bmatrix}$$

- Using mean-centred data U, compute the covariance matrix Σ

$$\Sigma = U^T U / (N - 1)$$

- We can then determine $\Phi^T \Sigma \Phi = \Lambda$
 - **Where the eigenvectors are in an orthogonal matrix - Φ**
 - We can consider Φ as a linear transformation
 - **And the eigenvalues are in an ordered diagonal matrix - Λ**
 - The amplitudes of the eigenvalues in Λ are proportional to the percentage of overall variance accounted for by the associated eigenvector

- **Facial recognition with PCA**

- **Common example application of PCA**
- Each face image is treated as a vector with $n = \text{rows} \times \text{columns}$ dimensions
- **Approach**
 - Create a matrix D
 - Where each face image is a row
 - D is an $N \times n$ matrix
 - Create mean centred data U
 - Determine the mean of the rows (face images)
 - Subtract from all the rows
 - Compute the covariance matrix
 - Solve for eigenvalues and eigenvectors
 - Normalise the eigenvectors
 - Select m Eigenvectors (with the largest m Eigenvalues)
 - Usually 20-50 for facial recognition
 - For an unknown image find its location in PCA space
 - Look for closest match

- **Performance**

- **Two aspects**
 - **Computation Time**
 - How long did it take?
 - **Success and Failure Rates**
 - What is the correct answer? We need **Ground Truth**
 - **Ground truths have to be manually computed - very difficult to get agreement**
 - How do we assess success? We need **Metrics**
 - **Compute TP TN FP FN**
 - **Then Compute:**

$$Recall = \frac{TP}{TP + FN}$$

$$Precision = \frac{TP}{TP + FP}$$

$$Accuracy = \frac{TP + TN}{Total\ Samples}$$

$$Specificity = \frac{TN}{FP + TN}$$

$$F_{\beta} = (1 + \beta^2) \cdot \frac{Precision \cdot Recall}{(\beta^2 \cdot Precision) + Recall}$$

- Can use these metrics to guide system tuning