

Video

- Video

- Aim is to reliably track moving objects of interest in a scene
 - Motion Detection
 - Moving object detection and location
 - Derivation of 3D object properties
- **What is an object of interest?**
 - Size - can only track reasonably large objects
 - max/min velocity and acceleration - can't go too fast, can't accelerate too quickly
 - Assumptions:
 - Mutual correspondence of object appearance
 - Common motion - e.g. walking, arms and legs move periodically

- Common Problems

- Illumination and appearance changes
 - Gradual (time of day), sudden (clouds/lights), shadows, weather (rain/snow)
- Background changes
 - Objects becoming part of the background
 - Objects leaving the background
 - Background objects oscillating slightly
- Setup
 - Camera motion
 - Frame rate
 - Field of view
 - Distance to objects
 - Location of camera

- Difference images

- Image subtraction

$$d(i,j) = |f_k(i,j) - b(i,j)|$$
$$d(i,j) = \begin{cases} 0 & \text{if } |f_k(i,j) - b(i,j)| < T \\ 1 & \text{otherwise} \end{cases}$$

- First formula retains greyscale, second is binary
- Difference between current frame and background model can highlight objects of interest
 - Colour images - Per channel difference? Just process hue?
- Won't always be right
 - False positives - incorrect object of interest identification
 - False negatives - incorrect background identification
 - Thresholding
 - Too high = not enough movement detected
 - Too low = too much accepted as objects of interest
 - **Dependent on high contrast between bg model and o.o.i**

- **Background Models**

- **Static background**

- Simplest approach
 - Sensitive to threshold T
 - How to get a background model?
 - First frame? Naive

- **Running average**

- $$b_{n+1}(i, j) = \alpha \cdot f_n(i, j) + (1 - \alpha) \cdot b_n(i, j)$$

- Where α is the learning rate
 - Adapts to a changing scene
 - **But...**
 - incorporates moving objects
 - Can try to avoid this by using great number of frames, or small learning rate, but makes slow to adapt

- **Median**

- Middle value from an ordered list

- $$h_n(i, j, p) = \sum_{k=(n-m+1) \dots n} \begin{cases} 1 & \dots \text{ if } (f_k(i, j) = p) \\ 0 & \dots \text{ otherwise} \end{cases}$$

- This creates a histogram
 - The median value for each pixel over m frames, assuming current frame is n
 - Median value is computed using this histogram
 - Every new frame that we see, the oldest value is discarded and the new one is added
 - **Very expensive** - store each frame being used to compute, as well as a histogram for each pixel
 - Determine a value for m and limit the histogram quantisation
 - **Update the histogram using aging (computationally inexpensive)**

- $$h_n(i, j, p) = \sum_{k=1 \dots n} \begin{cases} w_k & \dots \text{ if } (f_k(i, j) = p) \\ 0 & \dots \text{ otherwise} \end{cases}$$

where $w_1 = 1$ and $w_k = w_{k-1} * 1.001$

- Could do
 - Selective update
 - Mode
 - $b_n(i, j) = p$ where $h_n(i, j, p) \geq h_n(i, j, q)$ for all $q \neq p$

- **Gaussian Mixture Model**

- Approach to dealing with multi-modal background pixels
 - E.g. trees in wind, water etc
 - Model multiple values (3-5) at each point
 - Unsupervised learning
 - **Most popular methods for background modelling**

- Approach

• Model each pixel $f_n(i,j)$ using k Gaussian distributions each with

• $\pi_n(i,j,m)$

• $\mu_n(i,j,m)$

• $\sigma_n^2(i,j,m)$

- Weighting
- Mean
- Std. Deviation
- Set a learning constant α (where $0.01 \leq \alpha \leq 0.1$)
- For each new sample $f_n(i,j)$

• Select the best close Gaussian distribution

• close = within $2.5 \sigma_n(i,j,m)$ of $\mu_n(i,j,m)$

• If there is a best close Gaussian l

$$\pi_{n+1}(i,j,m) = \alpha * O_n(i,j,m) + (1 - \alpha) * \pi_n(i,j,m)$$

where $O_n(i,j,m) = 1$ for the close Gaussian distribution and 0 otherwise

$$\mu_{n+1}(i,j,m) = \mu_n(i,j,m) + O_n(i,j,m) * (\alpha / \pi_{n+1}(i,j,m)) * (f_n(i,j) - \mu_n(i,j,m))$$

$$\sigma_{n+1}^2(i,j,m) = \sigma_n^2(i,j,m) +$$

$$O_n(i,j,m) * (\alpha / \pi_{n+1}(i,j,m)) * ((f_n(i,j) - \mu_n(i,j,m))^2 - \sigma_n^2(i,j,m))$$

• If there is no close Gaussian (replace one...)

• $x = \operatorname{argmin}_m(\pi_n(i,j,m))$

• $\mu_{n+1}(i,j,x) = f_n(i,j)$

• $\sigma_{n+1}^2(i,j,x) = 2 * \max_m \sigma_n^2(i,j,m)$

- When there isn't a similar enough Gaussian, replace the smallest one
- **Identifying background distributions**
 - Define T , a proportion of frames in which the background pixels should be visible
 - Order the Gaussians by $\pi_{n+1}(i,j,m) / \sigma_{n+1}(i,j,m)$
 - Gaussians 1..B are considered background where

$$B = \operatorname{argmin}_b ((\sum_{b=1..m} \pi_{n+1}(i,j,m)) > T)$$

- Check if best close Gaussian (or the new Gaussian) is a background distribution
- **Use dilations and erosions to remove small regions and fill holes**
- **Issues with GMM**
 - Fails under fast variations
 - Low sensitivity to Gaussian tails
 - Less frequent events produce low probability & high variance
 - Needs to compute floating point probabilities

- **Codebook**

- ~~Designed to combat shortcomings of GMM~~
 - ~~Models each pixel independently~~
 - ~~For each pixel a codebook is maintained of RGB values (R_i, G_i, B_i)~~

- **Shadow detection**

- Compare current frame to background image
 - **Intensity/luminance decreases**
 - **Not by too much**
 - **Saturation doesn't increase too much**
 - **Neither does hue**
- Hue unpredictable and change in luminance can be small

$$SP_k(i, j) = \begin{cases} 1 \dots & \text{if } \left(\alpha < \frac{f_k^V(i, j)}{B_k^V(i, j)} < \beta \right) \text{ and } \left((f_k^S(i, j) - B_k^S(i, j)) < \tau_S \right) \text{ and } (|f_k^H(i, j) - B_k^H(i, j)| < \tau_H) \\ 0 \dots & \text{otherwise} \end{cases}$$

$$SP_k(i, j) = \begin{cases} 1 \dots & \text{if } \left(\lambda < \frac{f_k^V(i, j)}{B_k^V(i, j)} < 1.0 \right) \text{ and } (|f_k^S(i, j) - B_k^S(i, j)| < \tau_S) \\ 0 \dots & \text{otherwise} \end{cases}$$


- **Tracking**

- Used in surveillance, sports video analysis, car driving systems etc
- **Difficulties arise as objects:**
 - May be undergoing complex motion
 - May change shape
 - May be occluded
 - May change appearance due to lighting/weather
 - May physically change appearance
- **Exhaustive Search**
 - Extract object to be tracked from frame
 - Compare in all possible positions in future frame(s)
 - **Using a similarity metric e.g. normalised cross correlation**
 - Requires four degrees of freedom
 - Two for position
 - Change in x and y coordinates
 - One for scale
 - Movement towards camera (increase size), away from camera (decrease size)
 - One for rotation
 - **Can restrict our local maxima search using assumptions about the amount of change allowed between frames**
 - Will fail if motion is too complex
- **Mean Shift**
 - Track objects by
 - Searching locally for the most similar region
 - Using a histogram to represent the object
 - Using iterative gradient ascent/hill climbing to locate best match

- Histogram of colours, oriented gradients, textures etc.
- Weighting of centres is far greater than the edge colours
- **Modelling the object**
 - Probability density function (histogram) of colours

$$\hat{q}_u \triangleq C \sum_{i=1}^n k \left(\left\| \frac{\mathbf{x}_i}{h_q} \right\|^2 \right) \delta [b(\mathbf{x}_i) - u]$$

- Limiting the number of bins
- Typically use an elliptical region
- Weight the values relative to their location - **epanechnikov kernel**



- **Model candidate regions**

$$\hat{p}_u(\mathbf{y}) \triangleq C_h \sum_{i=1}^{n_h} k \left(\left\| \frac{\mathbf{y} - \mathbf{x}_i}{h_p} \right\|^2 \right) \delta [b(\mathbf{x}_i) - u]$$

- Matching to find the best new location
 - Compare histograms directly
 - Move to the mode in matching space
 - NCC observes the best results - clear spike

- **Mean Shift Approach**

- Consider the gradient of the similarity function (the Bhattacharyya coefficient)
- Gradient of superposition of kernels, centred at each data point is equivalent to convolving the data points with the gradient of the kernel

$$w_i = \sum_{u=1}^m \sqrt{\frac{\hat{q}_u}{\hat{p}_u(\mathbf{y}_0)}} \delta [b(\mathbf{x}_i) - u] \quad \mathbf{y}_1 = \frac{\sum_{i=1}^{n_h} \mathbf{x}_i w_i}{\sum_{i=1}^{n_h} w_i}$$

- Derived from the Bhattacharyya similarity measure
- Assumes Epanechnikov's kernels
- Moves in direction of highest gradient
- **Iterate until convergence - separation between \mathbf{y}_0 and $\mathbf{y}_1 < \text{a convergence threshold } (\epsilon)$**

- **Parameters**

- Number of bits poorer channel, convergence parameter (ϵ), kernel type

- **Background exclusion**

- If a background model is available...
 - Favour images which are similar to the object model && dissimilar to the corresponding background region
- **Multipart model**
 - Histograms lack spatial structure - can use a multipart model
 - Performance improves - **16 is optimal**
- **Dense Optical Flow**
 - Compute a motion field (known as an optical flow) for the entire image
 - **Direction and magnitude**
 - Based on the **brightness constancy constraint**
 - **Object points will have the same brightness over a short period of time**

$$f_t(i, j) = f_{t+\Delta t}(i + \Delta i, j + \Delta j)$$

- Need to find the displacement $(\Delta i, \Delta j)$ which will minimise the residual error

$$\varepsilon(\Delta i, \Delta j) = \sum_{i=i_{\text{current}}-w}^{i_{\text{current}}+w} \sum_{j=j_{\text{current}}-w}^{j_{\text{current}}+w} f_t(i, j) - f_{t+\Delta t}(i + \Delta i, j + \Delta j)$$

- To compute the optical flow $\left(\frac{\Delta i}{\Delta t}, \frac{\Delta j}{\Delta t}\right)$, assuming that displacement is small...

$$f_{t+\Delta t}(i + \Delta i, j + \Delta j) = f_t(i, j) + \frac{\partial f}{\partial i} \Delta i + \frac{\partial f}{\partial j} \Delta j + \frac{\partial f}{\partial t} \Delta t$$

- Hence...

$$\frac{\partial f}{\partial i} \Delta i + \frac{\partial f}{\partial j} \Delta j + \frac{\partial f}{\partial t} \Delta t = 0$$

So

$$\frac{\partial f}{\partial i} \frac{\Delta i}{\Delta t} + \frac{\partial f}{\partial j} \frac{\Delta j}{\Delta t} + \frac{\partial f}{\partial t} = 0$$

And reorganising

$$\begin{bmatrix} \frac{\partial f}{\partial i} & \frac{\partial f}{\partial j} \end{bmatrix} \begin{bmatrix} \frac{\Delta i}{\Delta t} \\ \frac{\Delta j}{\Delta t} \end{bmatrix} = - \frac{\partial f}{\partial t}$$

- **Problems**
 - Just shows the apparent motion of points, in a scene, based on brightness patterns
 - What happens when the brightness changes (brightness constancy does not hold)
 - **Perhaps look at optical flow in the gradient space**
 - What if a point moves differently to its neighbours
 - **Use regions based optical flow**

- What happens to a rotating sphere? Or a barber pole?
 - **It fails**
 - What happens if the motion is too large?
 - **Use iterative refinement**
- **Feature Based Tracking**
 - **Feature based optical flow**
 - We cannot compute optical flow for constant regions or along edges
 - **Better to compute just for features**