

## Parallel Programming in Haskell

- Parallelism has long been discussed in functional programming fields due to the immutable data and side-effect free computation
- No worry of deadlocks, starvations, synchronization, and locks
- Have to pay attention to execution granularity though

## Fibonacci Sequence calculator

### A small example

```
1 fib :: Integer -> Integer
2 fib n | n < 2 = 1
3 fib n = fib (n-1) + fib (n-2)
4
5 main = print $ fib 37
```

## Spark Parallelism

- Control.Parallel library
- This gives us the par operation, allowing us to signal sites of potential parallelism

```
par :: a -> b -> b
```

Semantically `par x y` is equivalent to just `y`, but the runtime is allowed to use it as a hint.

Next try

```
1 import Control.Parallel
2
3 fib :: Integer -> Integer
4 fib n | n < 2 = 1
5 fib n = par nf ( fib (n-1) + nf )
6         where nf = fib (n-2)
7
8 main = print $ fib 37
```

- **This makes the program slower.**
- Why?
  - `Par nf ( fib (n-1) + nf )`
  - In parallel evaluation, the RHS depends on the `nf` value, when the `nf` value is computed, a new `nf` dependency is introduced in the evaluation of `fib (n-1)`
  - **This is caused by the order in which (+) evaluates its arguments**

**We shouldn't need to consider something as low level as the order in which (+) handles its arguments...**

- `pseq`
- `pseq :: a -> b -> b`
- Evaluate `a` before `b`, returning `b`

```
1 fib n | n < 2 = 1
2 fib n = par nf1 (pseq nf2 (nf1 + nf2))
3         where nf1 = fib (n-1)
4               nf2 = fib (n-2)
```

- Evaluate `nf1` while at the same time (evaluate `nf2` and then return `nf1 + nf2`)

- **Much faster**

### Further Tweaking

- Overheads can dominate after a while
- Limit the number of new threads to allow more even distribution of work

```

1  import Control.Parallel
2
3  -- Sequential version of fib, when we want to avoid parallelism
4  sfib :: Integer -> Integer
5  sfib n | n < 2 = 1
6  sfib n = sfib (n-1) + sfib (n-2)
7
8  fib :: Integer -> Integer -> Integer
9  fib 0 n = sfib n
10 fib _ n | n < 2 = 1
11 fib d n = par nf1 (pseq nf2 (nf1 + nf2))
12           where nf1 = fib (d-1) (n-1)
13                 nf2 = fib (d-1) (n-2)
14
15 main = print $ fib 3 37

```

- **Parameterised thread count**
  - When thread count hits zero, evaluate remaining section of the numbers sequentially

### The Eval Monad

- Control.Parallel has more to it
- **Can separate algorithm and evaluation strategy**
  - **Eval Monad**
  - Semantically an Identity monad, in the same way that par is an identity function

```

5  runEval :: Eval a -> a
6
7  rpar :: a -> Eval a
8  rseq :: a -> Eval a

```

- rpar : My argument could be evaluated in parallel
- rseq : Evaluate my argument and wait for the result

```

1  fib d n = runEval $ do
2    nf1 <- rpar $ fib (d-1) (n-1)
3    nf2 <- rseq $ fib (d-1) (n-2)
4    return $ nf1 + nf2

```

- The **Control.Parallel.Strategies** module defines many utilities for easily exploiting parallelism
  - **Strategies for specifying strictness (play the role of seq)**
  - Higher level functions for applying these (**parList**, **parListChunk**, etc) which can model algorithms like map-reduce
- **Example of Strategy** - Parallelizing a Sudoku Solver
  - (Assuming the existence of a solver: solve :: String -> Maybe Grid)

```

1  import Sudoku
2  import Control.Exception
3  import System.Environment
4  import Data.Maybe
5
6  main :: IO ()
7  main = do
8      [f] <- getArgs
9      grids <- fmap lines $ readFile f
10     print (length (filter isJust (map solve grids)))

```

Version 1 is sequential:

```
$ ./sudoku-par1 sudoku17.1000.txt +RTS -N4 -s
```

```

...
Parallel GC work balance: 1.87% (serial 0%, perfect 100%)
...

```

Now let's add some basic parallelism

```

1  main :: IO ()
2  main = do
3      [f] <- getArgs
4      grids <- fmap lines $ readFile f
5
6      let (as,bs) = splitAt (length grids `div` 2) grids
7
8      solutions = runEval $ do
9          as' <- rpar (force (map solve as))
10         bs' <- rpar (force (map solve bs))
11         _ <- rseq as'
12         _ <- rseq bs'
13         return ( as' ++ bs' )
14     print (length (filter isJust solutions))

```

- Version 2 is far more efficient
- We can build Utilities in the Eval Monad

```

1  parallelMap :: (a -> b) -> [a] -> Eval [b]
2  parallelMap f [] = return []
3  parallelMap f (a:as) = do
4      b <- rpar (f a)
5      bs <- parallelMap f as
6      return (b:bs)
7
8  main :: IO ()
9  main = do
10     [f] <- getArgs
11     file <- readFile f
12
13     let puzzles = lines file
14         solutions = runEval (parallelMap solve puzzles)
15
16     print (length (filter isJust solutions))

```

- The library provides lots of pre-defined utilities:
  - Strategies: Control.Parallel.Strategies

- Evaluation strategies represent a parameterized hof to capture attempts to introduce parallelism
- **Type Strategy a = a -> Eval a**
- **Evaluate pairs in parallel**

```

1  parPair :: Strategy (a,b)
2
3  parPair (a,b) = do a' <- rpar a
4                    b' <- rpar b
5                    return (a',b')

```

- Nice way to apply these strategies: `using`

```

1  using :: a -> Strategy a -> a
2  x `using` s = runEval (s x)

```

- Strategies in the library are parameterised to allow us more control

```

1  evalPair :: Strategy a -> Strategy b -> Strategy (a,b)
2  evalPair sa sb (a,b) = do a' <- sa a
3                           b' <- sb b
4                           return (a',b')

```

which allows us to make various pairwise evaluation strategies:

```

1  parPair = evalPair rpar rpar
2  parSeqPair = evalPair rpar rseq

```

## Available library strategies

- `r0 :: Strategy a`
- `rseq :: Strategy a`
- `rpar :: Strategy a`
- `rparWith :: Strategy a -> Strategy a`
- `rdeepseq :: NFData a => Strategy a`

the `NFData` class represents data which can be fully evaluated.

```
1 rnf :: a -> ()
2 rnf a = a `seq` ()
```

- `rparWith :: Strategy a -> Strategy a`

This can “wrap” an `rpar` around it’s argument strategy. For example:

```
1 parPair sa sb = evalPair (rparWith sa) (rparWith sb)
```

Instead of the very parallel `parallelMap` we developed, the library has a parameterised implementation using some utilities:

```
1 evalList :: Strategy a -> Strategy [a]
2 evalList s [] = return []
3 evalList s (x:xs) = do x' <- s x
4                       xs' <- evalList s xs
5                       return (x':xs')
1 parList :: Strategy a -> Strategy [a]
2 parList s = evalList (rparWith s)
3
4 parMap :: Strategy b -> (a -> b) -> [a] -> [b]
5 parMap s f = (map f) `using` (parList s) . map f
```

Our Sudoku solver could have been written

```
1 let solutions = map solve puzzles `using` parList rseq
```

Many more utilities and strategies in the `Control.Parallel` module