

Classification

- Logistic Regression with Two Classes

- Two classes

- Output/target y is the **label** and takes a value of -1 or 1
- **Classifier predicts the label of a new object**
 - E.g. if an email is spam or not spam
- **Hypothesis**
 - Fitting a straight line to the data doesn't seem appropriate so not using same as lin reg

- **Predict 1 when $\theta^T x \geq 0$ and -1 when $\theta^T x < 0$**



- **So $h_{\theta}(x) = \text{sign}(\theta^T x)$**

- Logistic regression is essentially trying to fit a plane that separates $Y = 1$ data from the $Y = 0$ data

- $\theta^T x$ defines:

- **A point in one dimension** $1 + 0.5x_1 = 0 \rightarrow x_1 = -2$
- **A line in two dimensions** $2 + x_1 + 2x_2 = 0 \Rightarrow x_2 = \frac{-x_1}{2} - 1$
- ...
- **A plane in higher dimensions**

- Example: suppose x is vector $x = [1, x_1, x_2]^T$ e.g. x_1 might be tumour size and x_2 patient age.



- $\theta_0 = 0, \theta_1 = 0.5, \theta_2 = -0.5$.
- $h_{\theta}(x) \geq 0$ when $0.5x_1 - 0.5x_2 \geq 0$ i.e. when $x_1 \geq x_2$.
- When data can be separated in this way we say that it is "linearly separable".

- Not all data is linearly separable

- Choosing cost function

- Given training data

- $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$

$$x \in \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_n \end{bmatrix}, x_0 = 1, y \in \{-1, 1\}$$



- Create x_0 and set it to 1 so we can have compatible dimensions for $\theta^T x$
- And our hypothesis:
 - $h_\theta(x) = \text{sign}(\theta^T x)$
- How do we choose θ ? We **need a cost function**
 - **Maybe the 0-1 loss function**

$$\frac{1}{m} \sum_{i=1}^m \mathbb{I}(h_\theta(x^{(i)}) \neq y^{(i)})$$

- where indicator function $\mathbb{I} = 1$ if $h_\theta(x^{(i)}) \neq y^{(i)}$ and $\mathbb{I} = 0$ otherwise.
- **Essentially a count of the number of incorrect predictions, divided by the number of predictions**
- **Hard to work with**
- **For logistic regression we use:**

$$\frac{1}{m} \sum_{i=1}^m \log(1 + e^{-y^{(i)} \theta^T x^{(i)}}) / \log(2)$$

- **Noting:** $y = -1$ or $y = 1$
- Scaling by $\log(2)$ is optional **but makes loss = 1 when** $y^{(i)} \theta^T x^{(i)} = 0$



- Left is $y=1$, right is $y=-1$
- Large penalties occur when prediction is much less than 0 and $y = 1$ and when prediction is $\gg 0$ and $y = -1$
- Minimal when prediction almost equals label
- **So Logistic regression:**

Hypothesis: $h_{\theta}(x) = \text{sign}(\theta^T x)$

Parameters: θ

Cost Function: $J(\theta) = \frac{1}{m} \sum_{i=1}^m \log(1 + e^{-y^{(i)} \theta^T x^{(i)}})$

Goal: Select θ that minimises $J(\theta)$

- Gradient Descent

- General concept: Start with some θ , iterate to vector of θ which minimises $J(\theta)$
- Gradient descent approach:

Start with some θ

Repeat:

for $j=0$ to n { $\text{temp}j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$ }
for $j=0$ to n { $\theta_j := \text{temp}j$ }

- The gradient descent maths

For $J(\theta) = \frac{1}{m} \sum_{i=1}^m \log(1 + e^{-y^{(i)} \theta^T x^{(i)}})$:

- $\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{m} \sum_{i=1}^m -y^{(i)} x_j^{(i)} \frac{e^{-y^{(i)} \theta^T x^{(i)}}}{1 + e^{-y^{(i)} \theta^T x^{(i)}}}$
- (Remember $\frac{d \log(x)}{dx} = \frac{1}{x}$, $\frac{d \exp(x)}{dx} = \exp(x)$ and chain rule $\frac{df(z(x))}{dx} = \frac{df}{dz} \frac{dz}{dx}$)

So gradient descent algorithm is:

- Start with some θ
- Repeat:

for $j=0$ to n { $\text{temp}j := \theta_j + \frac{\alpha}{m} \sum_{i=1}^m y^{(i)} x_j^{(i)} \frac{e^{-y^{(i)} \theta^T x^{(i)}}}{1 + e^{-y^{(i)} \theta^T x^{(i)}}}$ }
for $j=0$ to n { $\theta_j := \text{temp}j$ }

- $J(\theta)$ is convex, has a single minimum. Iteration moves downhill until it reaches the minimum

- So using Logistic regression in practice
 - **Similar to linear regression**
 - **Select learning rate α to ensure convergence in reasonable time**
 - **Use cross-validation** to select what features to include/exclude from x
 - **Use regularisation to incorporate prior knowledge of constraints** on parameters θ e.g. add penalty on θ^2 by changing cost function to:

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \log(1 + e^{-y^{(i)} \theta^T x^{(i)}}) + \lambda \sum_{j=1}^n \theta_j^2$$

- Logistic Regression with Multiple Classes

- Examples
 - Filtering emails into folders, classifying weather
- Now our y output takes values 0,1,2...
 - 0 if it's sunny, 1 if it's cloudy...

- **Make it into a binary problem**

- Train a classifier $h_{\theta}^{(i)}(x)$ for each class i
- Predicts the probability that $y = i$
- Training data: re-label data as $y = -1$ when $y \neq i$ and $y = 1$ when $y = i$